

COMP 206 – Software Systems

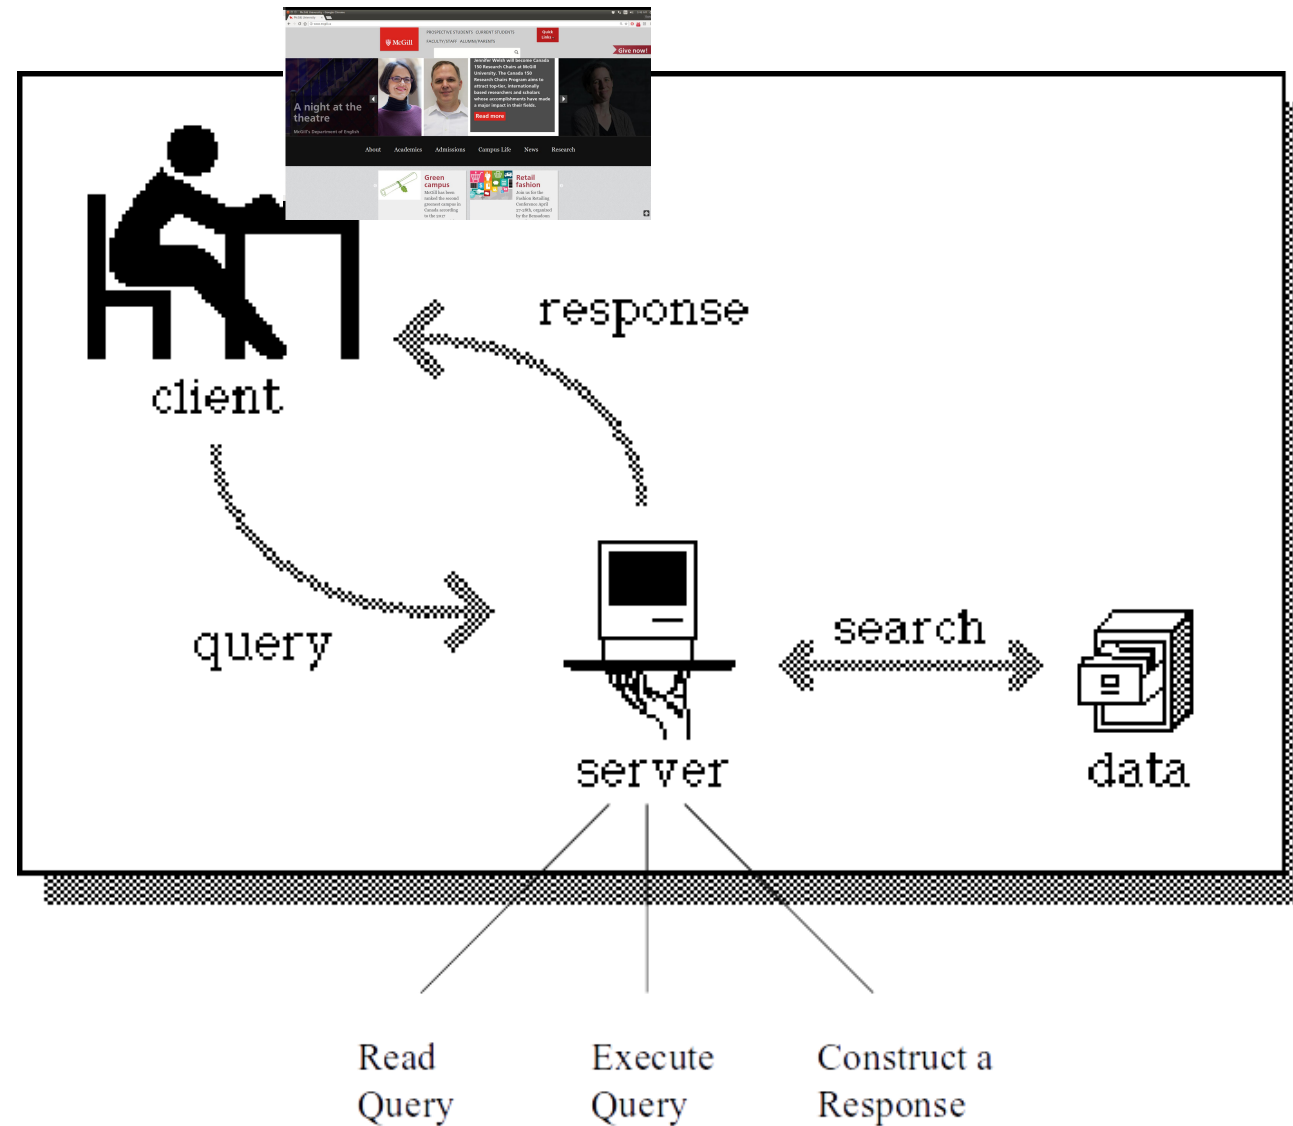
Nov 14th, 2018

Lecture 19 – The World Wide Web

Plan today

- Proof that 206 allows you to be a web programmer!
- The Web is an application layer on the Internet
 - It uses what we know about sockets, IP, TCP to move data between servers and browsers
 - Explanation of the protocols, connection to concepts we know well
 - Simple C code examples of a web browser and web server
- To achieve its goals, the Web defines a few new protocols:
 - HTTP (today)
 - HTML (today)
 - CGI (next lecture)
 - Many more that we don't cover here in 206

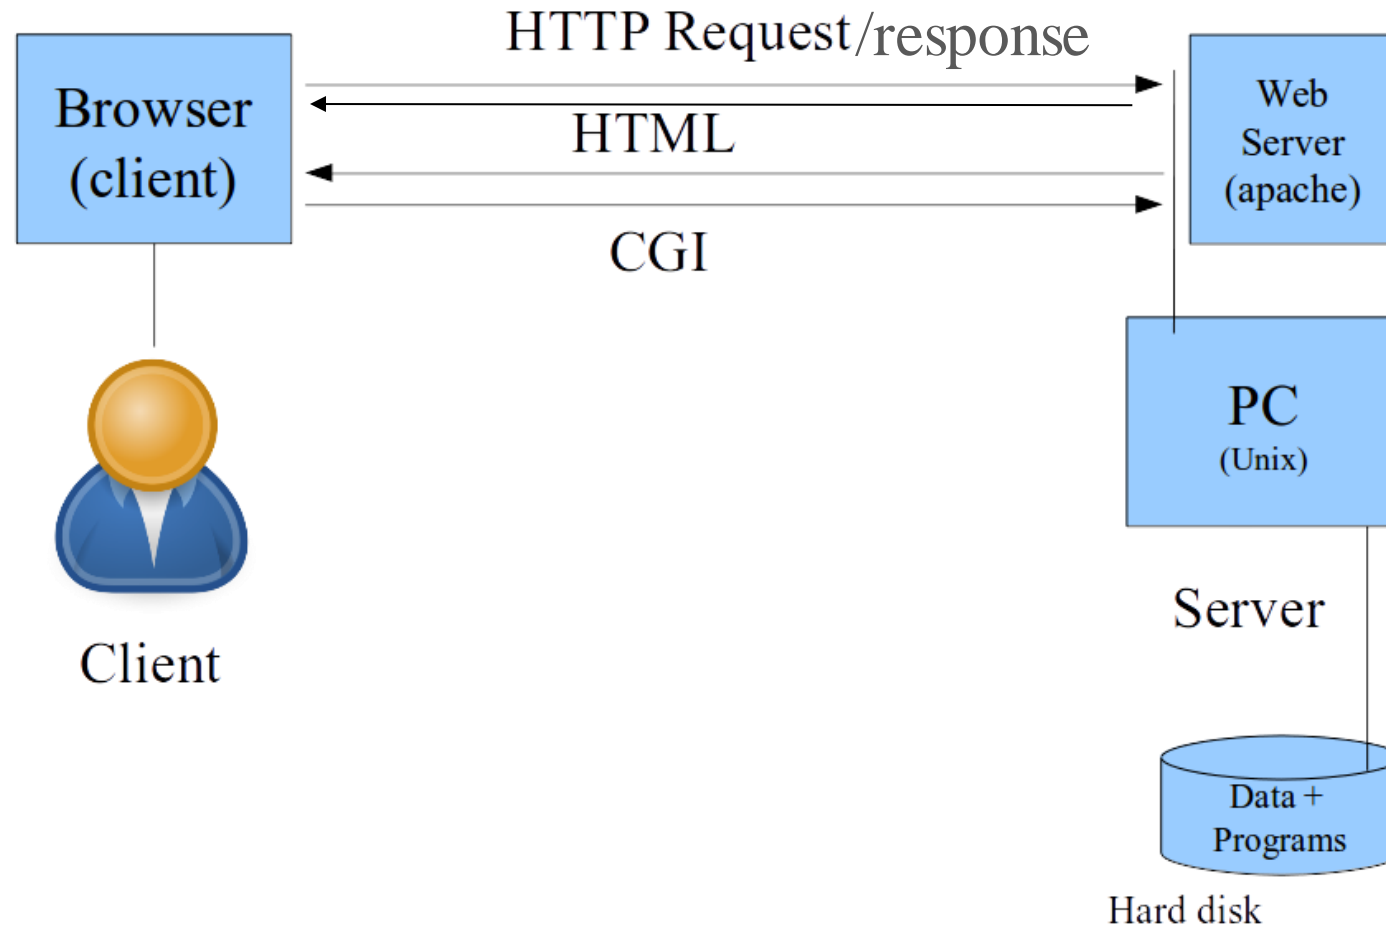
The Web as a Software System



World Wide Web

- The web is one of the most pervasive application-level software systems on the internet
- It's formed out of a collection of quite simple protocols that everyone can easily implement, so that many people, companies, governments, etc can interoperate even without using the same software:
 - Many browsers: chrome, safari, edge, lynx
- Let's start to understand those protocols and write some code to work with them

Web Software Protocols



Overview of Web Protocols

- HyperText Transmission Protocol (HTTP) : the backbone for web inter-connections. Describes how to request and interpret basic web data
- Hyper Text Markup Language (HTML) : the basic building-block of web pages. Text, formatting, colors, images, extensible components for new ideas (3D graphics, UIs, etc)
- Common Gateway Interface (CGI) : A protocol to allow server-side programs to be run through web interactions, producing output that drives subsequent interactions
- Many more elements now exist on the client side: javascript, css, flash, mobile application eco-systems, and many more (not for 206)

HyperText Transfer Protocol (HTTP)

- TCP/IP allows end-points to transfer arbitrary binary data streams.
- This is like giving the computer a mouth and ears, but without a language, how can we write code that accomplishes anything?
- HTTP is the "language" of web connections. It lets browsers ask for web content and lets servers provide it. It is an open, evolving, community-defined standard that all web entities use to structure their transmissions so that each side can understand:
 - <https://tools.ietf.org/html/rfc7230#section-2.1>

HTTP and Networking

- HTTP's job is to allow clients to request the resources they desire and to allow servers to respond appropriately.
- It lives on top of lower network layers. The request and response are plain-text payloads with a specific format that must be transmitted somehow (of course we all use TCP/IP, through sockets!)
- A HTTP client establishes a connection over a predefined port :
 - 80 for normal HTTP
 - 443 for SSL HTTP (secure)
- An HTTP server sits waiting on the port, parses the client request and responds (perhaps on a different port, but let's ignore that for now)

HTTP Requests and Responses

- Client starts the communication by sending an HTTP request that includes a ***METHOD, Uniform Resource Identifier (URI)***, protocol version, and header:
 - The most common method is GET, this means "Send me back the page"
- The server sends back a ***response code***, response text, header, followed by the requested data, if any. Common response codes:
 - 200: OK
 - 401: Unauthorized
 - 403: Forbidden
 - 404: Not Found
 - 500: Internal Server Error

URI vs URL

- We are used to typing Uniform Resource Locators (URLs) into our web browser. They include:
 - the server name/address,
 - optionally, the port if we care to switch away from the default 80/443 (separate by colon, like mimi.cs.mcgill.ca:5000)
 - and then describes the path to a file on that server.
- The URI can be the portion only ***after the server address and port***, while the URL must include all of the above
 - It makes sense only to need to URI within an HTTP request, because we are sending this data to a server on a port (full URL re-dundant)
 - Unless the server is having an identity crisis (midterms are hard ok!)

HTTP Session Example

Client request:

```
GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi
```

Server response:


```
HTTP/1.1 200 OK
Date: Mon, 27 Jul 2009 12:28:53 GMT
Server: Apache
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
ETag: "34aa387-d-1568eb00"
Accept-Ranges: bytes
Content-Length: 51
Vary: Accept-Encoding
Content-Type: text/plain

Hello World! My payload includes a trailing CRLF.
```

HTTP Session Example

METHOD,
Uniform Resource Identifier (URI),
protocol version
header

Client request:



```
GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi
```

Server response:

```
HTTP/1.1 200 OK
Date: Mon, 27 Jul 2009 12:28:53 GMT
Server: Apache
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
ETag: "34aa387-d-1568eb00"
Accept-Ranges: bytes
Content-Length: 51
Vary: Accept-Encoding
Content-Type: text/plain

Hello World! My payload includes a trailing CRLF.
```

HTTP Session Example

Client request:

METHOD,
Uniform Resource Identifier (URI),
protocol version
header

→ GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi

Server response:

HTTP/1.1 200 OK
Date: Mon, 27 Jul 2009 12:28:53 GMT
Server: Apache
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
ETag: "34aa387-d-1568eb00"
Accept-Ranges: bytes
Content-Length: 51
Vary: Accept-Encoding
Content-Type: text/plain

Hello World! My payload includes a trailing CRLF.

HTTP Session Example

***METHOD,
Uniform Resource Identifier (URI),
protocol version
header***

Client request:

GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi

Server response:


HTTP/1.1 200 OK
Date: Mon, 27 Jul 2009 12:28:53 GMT
Server: Apache
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
ETag: "34aa387-d-1568eb00"
Accept-Ranges: bytes
Content-Length: 51
Vary: Accept-Encoding
Content-Type: text/plain

Hello World! My payload includes a trailing CRLF.

HTTP Session Example

METHOD,
Uniform Resource Identifier (URI),
protocol version
Headers (everything else)

Client request:



```
GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi
```

Server response:

```
HTTP/1.1 200 OK
Date: Mon, 27 Jul 2009 12:28:53 GMT
Server: Apache
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
ETag: "34aa387-d-1568eb00"
Accept-Ranges: bytes
Content-Length: 51
Vary: Accept-Encoding
Content-Type: text/plain


Hello World! My payload includes a trailing CRLF.
```

HTTP Session Example

Client request:

```
GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi
```

Server response:

<i>response code</i>		HTTP/1.1 200 OK
response text		Date: Mon, 27 Jul 2009 12:28:53 GMT
header,		Server: Apache
the requested data		Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
		ETag: "34aa387-d-1568eb00"
		Accept-Ranges: bytes
		Content-Length: 51
		Vary: Accept-Encoding
		Content-Type: text/plain
		 Hello World! My payload includes a trailing CRLF.

HTTP Session Example

Client request:

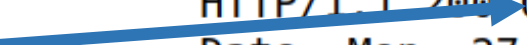
```
GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi
```

Server response:

response code
response text
header,
the requested data

```
HTTP/1.1 200 OK
Date: Mon, 27 Jul 2009 12:28:53 GMT
Server: Apache
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
ETag: "34aa387-d-1568eb00"
Accept-Ranges: bytes
Content-Length: 51
Vary: Accept-Encoding
Content-Type: text/plain

Hello World! My payload includes a trailing CRLF.
```




HTTP Session Example

Client request:

```
GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi
```

Server response:

response code
response text
header,
the requested data



```
HTTP/1.1 200 OK
Date: Mon, 27 Jul 2009 12:28:53 GMT
Server: Apache
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
ETag: "34aa387-d-1568eb00"
Accept-Ranges: bytes
Content-Length: 51
Vary: Accept-Encoding
Content-Type: text/plain

Hello World! My payload includes a trailing CRLF.
```

HTTP Session Example

Client request:

```
GET /hello.txt HTTP/1.1
User-Agent: curl/7.16.3 libcurl/7.16.3 OpenSSL/0.9.7l zlib/1.2.3
Host: www.example.com
Accept-Language: en, mi
```

Server response:

response code
response text
header,
the requested data

```
HTTP/1.1 200 OK
Date: Mon, 27 Jul 2009 12:28:53 GMT
Server: Apache
Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT
ETag: "34aa387-d-1568eb00"
Accept-Ranges: bytes
Content-Length: 51
Vary: Accept-Encoding
Content-Type: text/plain


Hello World! My payload includes a trailing CRLF.
```

Hyper Text Markup Language (HTML)

- An example of formatted text (we know it well from A2 Q2)
- An HTML document is made up of tags with content. Examples:
 - `<html></html>` encloses the full document
 - `<body></body>` encloses the main content of the page
 - `<a>` encloses information about a link
- Each tag can also have arguments. Example:
 - ``
- HTML is a convention that allows many web authors and many web browsers to see the web in a consistent way.
- It is not actually connected much with HTTP and the code that runs the web, so it's not a focus for 206 (although we will mess with it a bit when we write CGI functions)

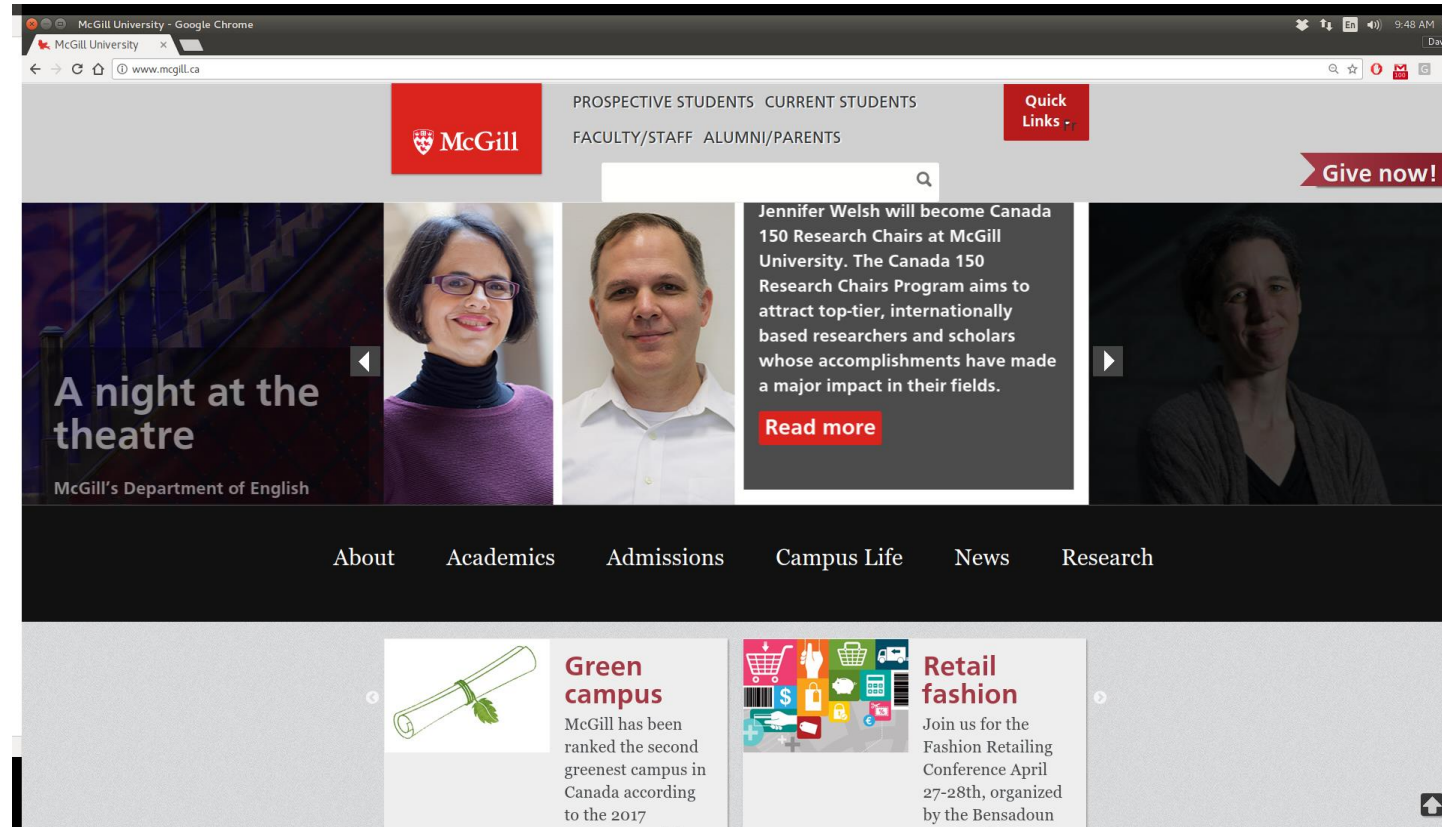
Your Web Browser as a Software System

- You see rendered HTML including images, UI elements, color. How did this get there?



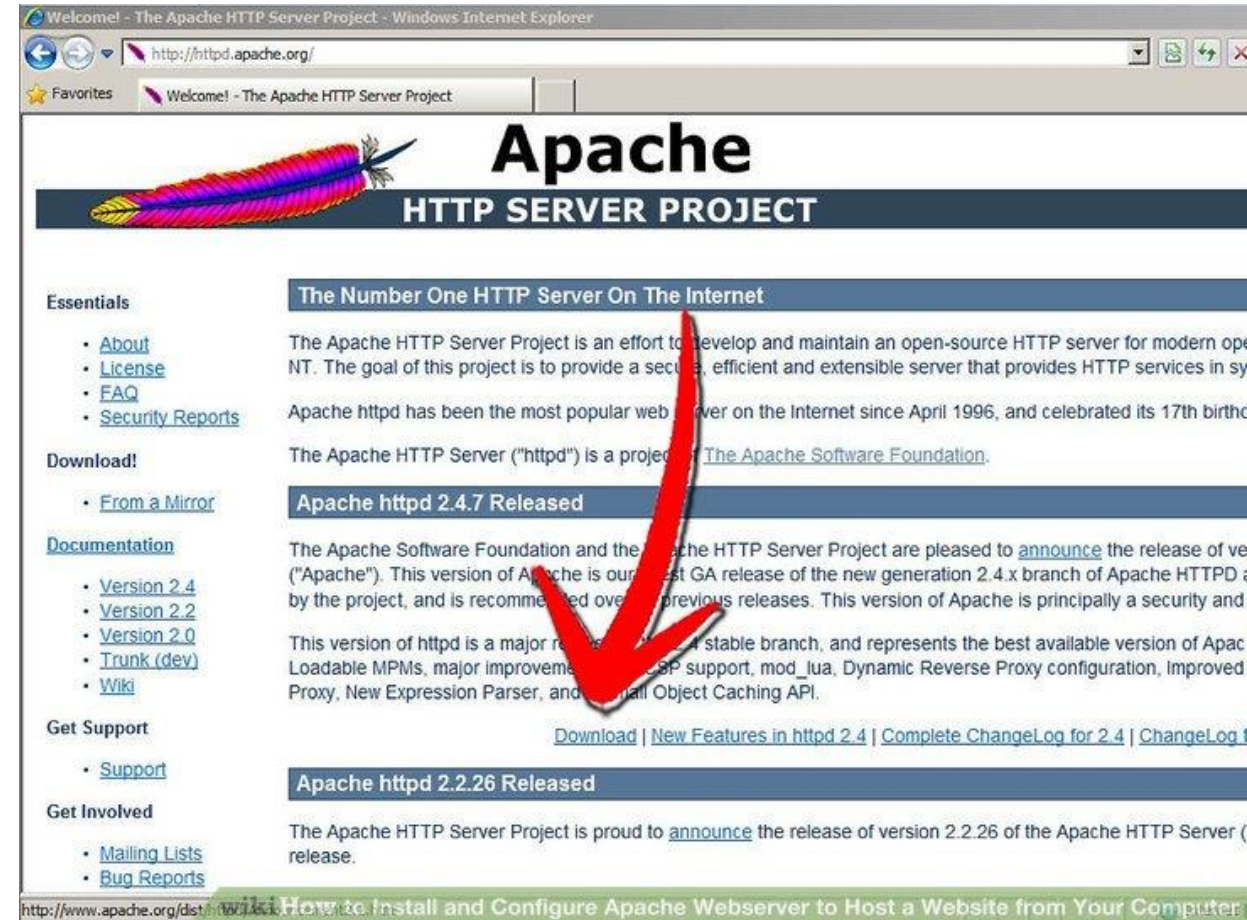
Your Web Browser as a Software System

- You see rendered HTML including images, UI elements, color. How did this get there?
- Your browser is making HTTP Requests and reading the Responses:
 - Started each time you enter a URL or click a link
 - Parses the response (HTML page, image, etc) and displays what's appropriate



The other side: web server software

- A website is accessible because some computer is running persistent server code waiting for browsers to connect
- The server holds the data (HTML pages, images, programs, etc)
- Upon a connection, searches for the most appropriate response and returns the payload



Managing HTTP in C

- Both the client (web browser) and server (web server) can be implemented in C quite easily with what we know now!
 - We will see a bit of this in Assignment 4
- Let's think for the next few slides about what the C code looks like and how we'd run it. Take a look in the ExampleCode folder on our Github to follow along

A Web Browser in C: pseudo-code

- Required to initiate the socket communications: `socket()` and `connect()`
 - Provide the server's address and use port number 80
- Required to form the HTTP Request. This is just a simple C string:
 - `char *request = "GET / HTTP1.1";` (if we want to get the index.html)
- Send with `write()` onto the socket
- Wait for the response and gather it using `read()` from the socket
- Display information received if the code was OK, else tell the user the error

Web Client Example

```
8 #include <stdio.h>
9 #include <stdlib.h>
10 #include <sys/socket.h>
11 #include <string.h>
12 #include <sys/types.h>
13 #include <netinet/in.h>
14 #include <netdb.h>
15
16 int main(int argc, char** argv) {
17
18     struct addrinfo hints;
19     memset(&hints, 0, sizeof hints);
20     hints.ai_family = AF_INET ;
21     hints.ai_socktype = SOCK_STREAM;
22     struct addrinfo *servinfo;
23     int status = getaddrinfo(argv[1], "80",
24                             &hints, &servinfo);
25     int sockfd = socket(servinfo->ai_family,
26                         servinfo->ai_socktype,
27                         servinfo->ai_protocol);
28     connect(sockfd,
29             servinfo->ai_addr,
30             servinfo->ai_addrlen);
31
32     char header[1000];
33     sprintf(header, "GET /index.html HTTP/1.1\r\nHost:%s\r\n\r\n", argv[1] );
34     int n = write(sockfd, header, strlen(header));
35
36     char buffer[2048];
37     n = read(sockfd, buffer, 2048);
38     printf("%s", buffer);
39
40     return (EXIT_SUCCESS);
41 }
42
```

Posted to Lecture19 folder
in ExampleCode

A Web Server in C: pseudo-code

- Loop forever awaiting connections, when one arrives:
 - Read the HTTP request, depending on the command and URI:
 - Perhaps look up a page on the local disk, read it into memory
 - Perhaps execute a CGI by forming the correct stdin, capture its stdout in memory (this to be covered in the next set of slides!)
 - Perhaps prepare an error response as a nicely formatted HTML page in memory
 - Write the HTTP response back to the client, forming the correct header and placing whatever we assembled in memory into the body

Web Server Example

- Ends up being too long to see on one slide reasonably. Check the file "httpd.c" in the Lecture19 folder in ExmampleCode

Exercises

- Run the browser and server yourself in a couple of places. You might get some problems with ports not being open, so be creative and use some trouble shooting (linux commands etc)
- Try to modify the provided browser and server example to change the HTTP requests in some way. Run and see what errors you get.
- Write an HTML page if you never have before. You can view it using your own browser using <file:///<path>> or by placing it in the htdocs folder and running our sample server
- Writing a web client yourself from scratch. This will ensure you "get it" with sockets and HTTP request/responses