

# lecture 2

- fixed point
- IEEE floating point standard

# Fixed point

Fixed point means we have a constant number of bits (or digits) to the left and right of the binary (or decimal) point.

Examples :

23953223.49 (base 10)

Currency uses a fixed number of digits to the right.

10.1101 (base 2)

# Two's complement for fixed point numbers

e.g. 0110.1000 which is 6.5 in decimal

**How do we represent -6.5 in fixed point ?**

$$\begin{array}{rcl} 0110.1000 & & \\ 1001.0111 & \leftarrow & \text{invert bits} \\ + 0000.0001 & \leftarrow & \text{add .0001} \\ \hline 0000.0000 & & \end{array}$$

Thus,

$$\begin{array}{rcl} 1001.0111 & \leftarrow & \text{invert bits} \\ + 0000.0001 & \leftarrow & \text{add .0001} \\ \hline 1001.1000 & \leftarrow & \text{answer: -6.5 in (signed) fixed point} \end{array}$$

# Scientific Notation (floating point)

$$300,000,000 = 3 \times 10^8$$
$$= 3.0 E + 8$$

$$.00000456 = 4.56 E - 6$$



"Normalized" : one digit to the left of the decimal point.

# Scientific Notation in binary

$$(1000.01)_2 = 1.00001 \times 2^3$$

$$(0.111)_2 = 1.11 \times 2^{-1}$$

"Normalized" means one "1" bit to the left of the binary point. **(Note that 0 cannot be represented this way.)**

sign

"exponent"

+

|

.

"significand"

(also called  
"mantissa")

x

2



<sup>E</sup>

How to represent this information ?

How to represent the number 0 ?

# IEEE floating point standard (est. 1985)

case 1: single precision (32 bits = 4 bytes)



Let's look at these three parts, and then examples.

sign                      0 for positive, 1 for negative

"significand"



You don't encode the "1" to the left of the binary point.

Only encode the first 23 bits to the right of the binary point.



exponent code

exponent value

00000000

reserved (explained soon)

00000001

-126

00000010

-125

00000011

- 124

:

:

:

:

01111111

0

This is not two's  
complement !

10000000

1

10000001

2

:

:

:

:

11111110

127

11111111

reserved (explained soon)

unsigned exponent code = exponent value + "bias"  
(for 8 bits, bias is defined to be 127)

Q: What is the largest positive normalized number ?  
(single precision)

$$1.\text{ } \times 2^e$$

A:

$$1.\text{ } 1111 \dots 11 \times 2^{127}$$

$$2^{127} \approx 10^?$$

$$2^{10} \approx 10^3$$

$$2^{127} = 2^{120} \cdot 2^7$$

$$= (2^{10})^{12} \cdot 2^7$$

$$\approx (10^3)^{12} \cdot 10^2$$

$$= 10^{38}$$

Q: What is the smallest positive normalized number ?  
(single precision)

$$1 \text{ . } \text{-----} \times 2^e$$

A:

$$1 \text{ . } 000000 \dots 0 \times 2^{-126}$$


Exponent code 00000000 reserved for  
"denormalized" numbers

$$1 \text{ + } \bigcirc . \text{ } \xrightarrow{\hspace{10cm}} \times 2^{-126}$$

• belong to  $(-2^{-126}, 2^{-126})$

• includes 0

Dividing each power of 2 interval into  $2^{23}$  equal parts  
(same for negative real numbers).



Note the power of 2 intervals themselves are equally spaced on a log scale.



Exponent code 11111111 also reserved.

if significand is all 0's

then value is  $\pm$  infinity (depending on sign bit)

else value is NaN ("not a number")

e.g. variable is declared but hasn't been  
assigned a value

Example: write 8.75 a single precision float (IEEE).

First convert to binary.

$$8.75$$

$$= (1000)_2 \cdot (75)_{10}$$

$$= (10001)_2 \cdot (5)_{10} \times 2^{-1}$$

$$= 100011.0 \times 2^{-2}$$

$$= 1.00011 \times 2^3$$



$$(8.75)_{10} = (1.00011)_2 \times 2^3$$

23 bit significand: 000110000000000000000000000

exponent value:  $e = 3$

exponent code = exponent value (e) + bias

Thus, exponent code is unsigned  $3 + 127$ .

$$(130)_{10} = (10000010)_2$$

So, the 32 bit representation is :

$\underbrace{0}_{0 \times 4} \underbrace{10000010}_1 \underbrace{0001}_{0} \underbrace{100000000000000000000000}_{c} \underbrace{0}_{0} \underbrace{0}_{0} \underbrace{0}_{0} \underbrace{0}_{0}$

Recall last lecture: 0.05 cannot be represented exactly.

```
float x = 0;  
for (int ct = 0; ct < 20; ct ++) {  
    x += 1.0 / 20;  
    System.out.println( x );  
}
```

0.05

0.1

0.15

0.2

0.25

0.3

0.35000002

0.40000004

0.45000005

0.50000006

**etc**

# Floating Point Addition

$$x = 1.00100100010000010100001 \quad * \quad 2^2$$

$$y = 1.10101000000000000101010 \quad * \quad 2^{-3}$$

$$x + y = ?$$

# Floating Point Addition

$$x = 1.00100100010000010100001 \quad * \quad 2^2$$

$$y = 1.10101000000000000101010 \quad * \quad 2^{-3}$$

$$x + y = ?$$

$$x = 1.0010010001000001010000100000 \quad * \quad 2^2$$

$$y = .000011010100000000000101010 \quad * \quad 2^2$$

but the result  $x+y$  has more than 23 bits of significand

How many *digits* (base 10) of precision can we represent with 23 *bits* (base 2) ?



$$2^{23}$$

$$= (2^{10})^2 \cdot 2^3 = (1024)^2 \cdot 2^3$$

$$\approx 10^6 \cdot 10^1$$

$$= 10^7$$

case 2: double precision (64 bits = 8 bytes)



## exponent code

## exponent value

unsigned exponent code = exponent value + bias

*For 11 bits, bias is defined to be  $2^{10} - 1 = 1023$ .*

000000000000

000000000001

000000000010

000000000011

⋮

011111111111

100000000000

100000000001

⋮

111111111110

111111111111

reserved

-1022

-1021

- 1020

⋮

0

1

2

⋮

1023

reserved

# Example

$$(8.75)_{10} = (1.00011)_2 \times 2^3$$

significand (52 bits)

$$= .0001\overline{1} = \frac{1}{10^4} = 0.0001$$

exponent = 3, code using 11 bits:

$$3 + 1023 = 1026 = (10000000010)_2$$

# double precision float (64 bits)

[illegible]

0 x 4      0      2      1      8      0      0      0      0      0      000000



Q: What is the largest positive normalized number ?  
(double precision)

$$1.1111111111111111 \times 2^e$$

A:

$$2^{1023}$$

$$11 \quad (2^{10})^{102} \quad 2^3$$

$$22 \quad (10^3)^{102} \quad 10$$

$$11 \quad 10^{307}$$

# Approximation Errors (Java/C/...)

```
double x = 0;  
for (int ct=0; ct < 10; ct++) {  
    x += 1.0 / 10;  
    System.out.println( x );  
}
```

0.1

0.2

0.30000000000000000004

0.4

0.5

0.6

0.7

0.7999999999999999999

0.8999999999999999999

0.9999999999999999999

How many *digits* of precision can we represent with 52 *bits* ?

$$\begin{aligned} & 2^{52} \\ = & (2^{10})^5 2^2 \\ \approx & (10^3)^5 10 \\ = & 10^{16} \end{aligned}$$

52 bits covers about the same "range" as 16 digits.

That is why the print out on the previous slide had up to (about) 16 digits to the right of the decimal point.