

Applied Machine Learning

Convolutional Neural Networks

Siamak Ravankhah

COMP 551 (winter 2020)

Learning objectives

understand the convolution layer and the architecture of conv-net

- its inductive bias
- its derivation from fully connected layer
- different types of convolution

MLP and image data

5	0	4	1
3	5	3	6
4	0	9	1
3	8	6	9
		↓	

we can apply an MLP to image data

first vectorize the input $x \rightarrow \text{vec}(x) \in \mathbb{R}^{784}$

feed it to the MLP (with L layers) and predict the labels

$$\text{softmax} \circ W^{\{L\}} \circ \dots \circ \text{ReLU} \circ W^{\{1\}} \text{vect}(x)$$

the model knows nothing about the image structure ***spatial info is NOT used.***

we could shuffle all pixels and learn an MLP with similar performance

how to bias the model, so that it "knows" its input is image?

image is like 2D version of sequence data

Let's find the right model for sequence first...

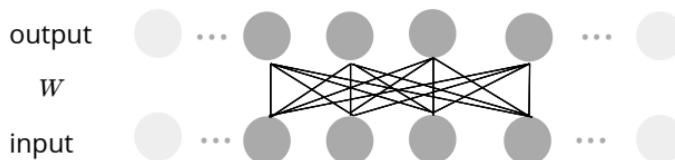
image:<https://medium.com/@raijatiajn0807/machine-learning-6ecde3bfd2f4>

Parameter-sharing

suppose we want to convert one sequence to another $\mathbb{R}^D \rightarrow \mathbb{R}^D$

suppose we have a dataset of input-output pairs $\{(x^{(n)}, y^{(n)})\}_n$

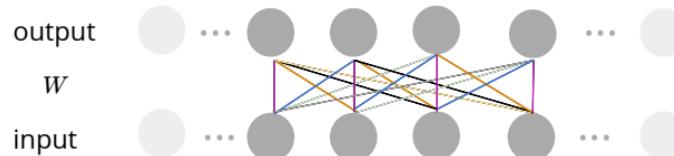
consider only a single layer $y = g(Wx)$



we may assume, each output unit is the same function shifted along the sequence

Assume input doesn't change much over time -

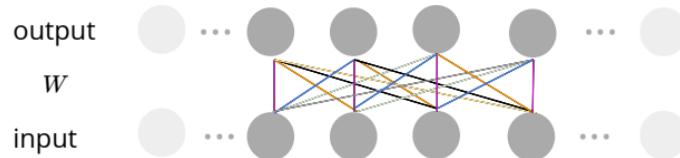
when is this a good assumption?



elements of w of the same color are tied together
(parameter-sharing)

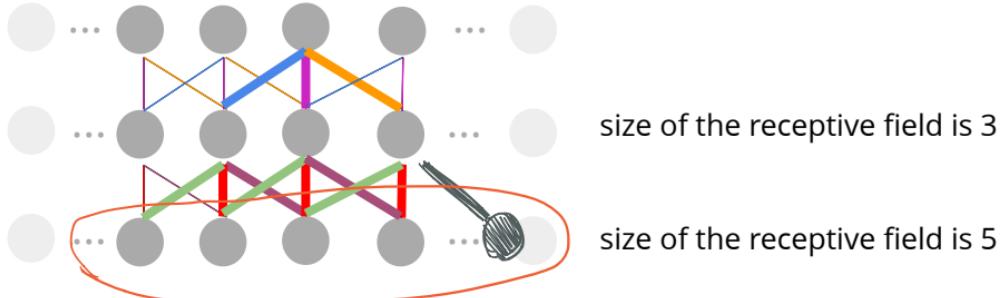
Locality & sparse weight

we may assume, each output unit is the same function shifted along the sequence



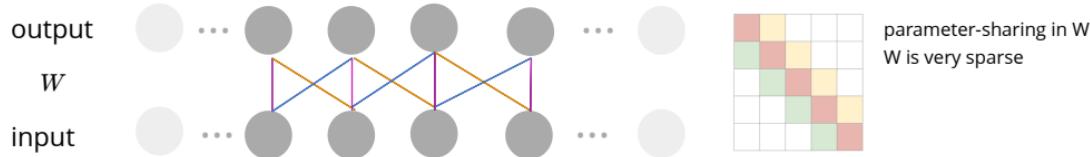
we may further assume each output is a **local** function of input

larger **receptive field** with multiple layers



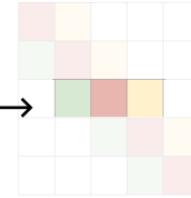
Cross-correlation (1D)

we may further assume each output is a local function of input



instead of the whole matrix we can keep the one set of nonzero values

$$w = [w_1, \dots, w_K] = [W_{c,c-\lfloor \frac{K}{2} \rfloor}, \dots, W_{c,c+\lfloor \frac{K}{2} \rfloor}] \longrightarrow$$



we can write matrix multiplication as **cross-correlation** of w and x

$$y_c = g\left(\sum_{d=1}^D W_{c,d} x_d\right) = g\left(\sum_{k=1}^K w_k x_{c-\lfloor \frac{K}{2} \rfloor + k}\right)$$

slide on the input, calculate inner product and apply the nonlinearity

Convolution (1D)

Cross-correlation is similar to convolution

Cross-correlation $y(c) = \sum_{k=-\infty}^{\infty} w(k)x(c+k)$

w is called the filter or kernel

ignoring the activation (for simpler notation)

assuming w and x are zero for any index outside the input and filter bound

3 4 5
1 0 2
-1 9 7
 \downarrow
7 2 1
9 0 4
-1 1 3

Convolution flips w or x (to be **commutative**)

$$y(c) = \sum_{k=-\infty}^{\infty} w(k)x(c-k) = \sum_{d=-\infty}^{\infty} w(c-d)x(d)$$

flip horizontally $w * x$ change of variable $x * w$
and vertically

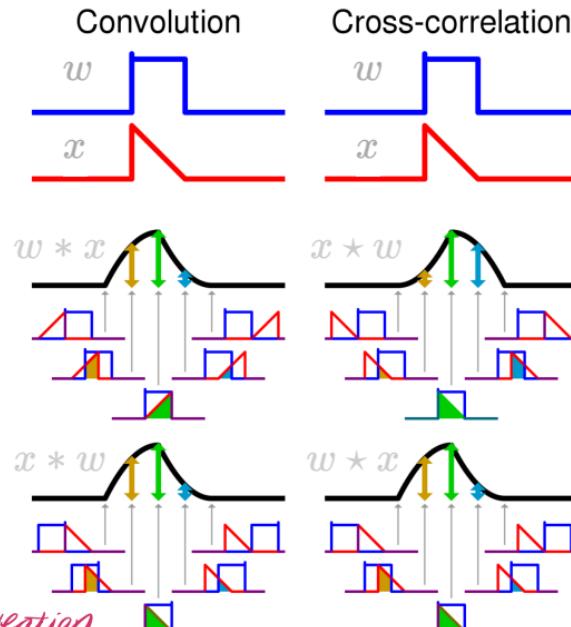
since we learn w, flipping it makes no difference

in practice, we use cross correlation rather than convolution

convolution is **equivariant** wrt translation

-- i.e., shifting x, shifts $w * x$

*but still call it
convolution by convention*



Convolution (2D)

similar idea of parameter-sharing and locality extends to 2 dimension (*i.e. image data*)

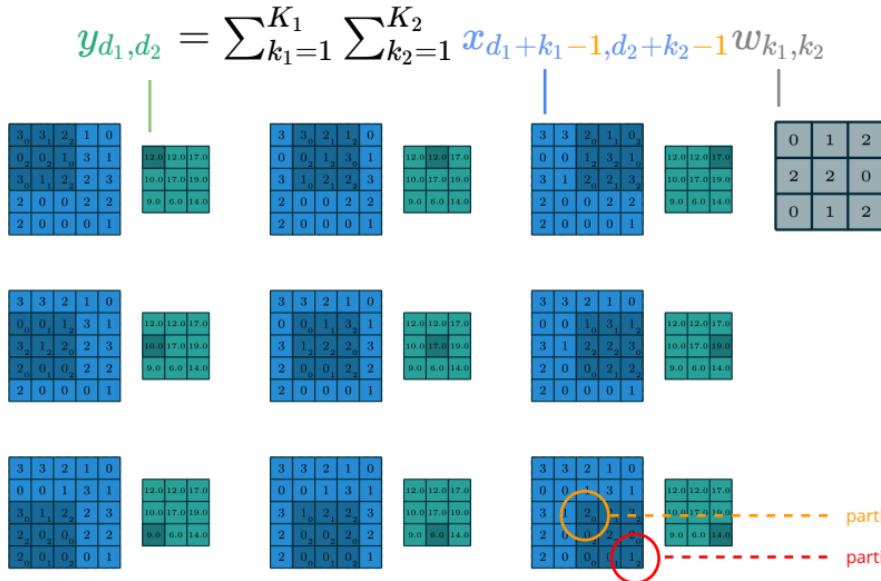


image credit: Vincent Dumoulin, Francesco Visin

Convolution (2D)

similar idea of parameter-sharing and locality extends to 2 dimension (*i.e. image data*)

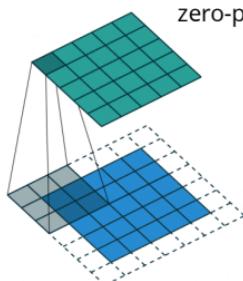
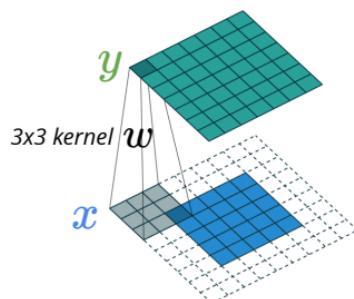
$$y_{d_1, d_2} = \sum_{k_1=1}^{K_1} \sum_{k_2=1}^{K_2} x_{d_1+k_1-1, d_2+k_2-1} w_{k_1, k_2}$$

there are different ways of handling the borders

zero-pad the input, and produce all non-zero outputs (**full**)

output is larger than input (by how much?)

each input participates in the same number of output elements



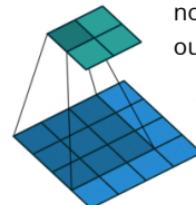
$$P = \frac{t-1}{s}$$

k is usually odd

output length (for one dimension)
 $\lfloor D + \text{padding} - K + 1 \rfloor$

\downarrow
 $2P$

zero-pad the input, to keep the output dims similar to input (**same**)



no padding at all (**valid**)

output is smaller than input (how much?)

- ① shrinking input (problematic!)
- ② throw the information at corner

image credit: Vincent Dumoulin, Francesco Visin

Pooling

sometimes we would like to reduce the size of output e.g., from $D \times D$ to $D/2 \times D/2$

a combination of pooling and downsampling is used

1. calculate the output $\tilde{y}_d = \text{green}(\sum_{k=1}^K x_{d+k-1} w_k)$

2. aggregate the output over different regions

$$y_d = \text{pool}\{\tilde{y}_d, \dots, \tilde{y}_{d+p}\}$$

two common aggregation functions are **max** and **mean**
pooling results in some degree of invariance to translation

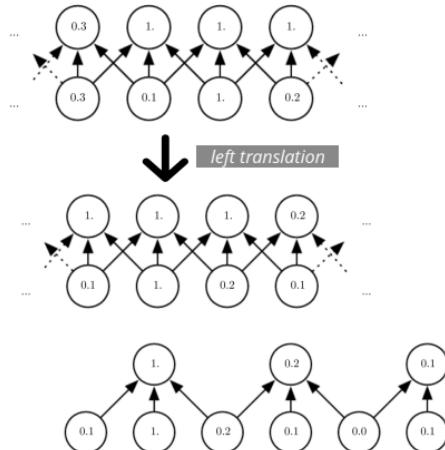
3. often this is followed by subsampling using the same step size

the same idea extends to higher dimensions

12	20	30	0
8	12	2	0
34	70	37	4
112	100	25	12

$\xrightarrow{2 \times 2 \text{ Max-Pool}}$

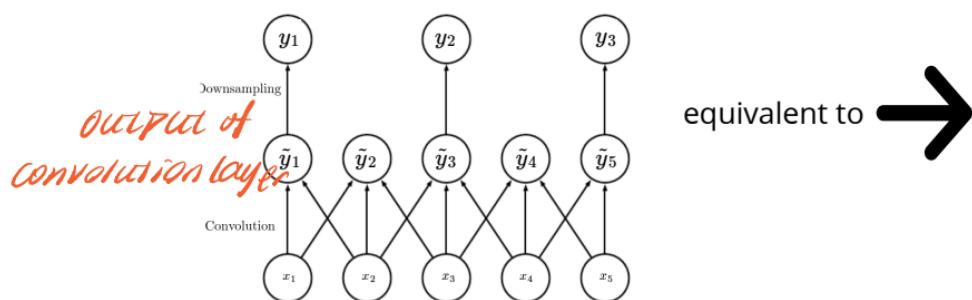
20	30
112	37



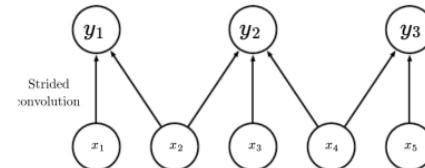
Strided convolution

alternatively we can directly subsample the output

$$\begin{aligned}\tilde{y}_d &= g\left(\sum_{k=1}^K x_{(d-1)+k} w_k\right) \\ y_d &= \tilde{y}_{dp}\end{aligned}$$



$$\tilde{y}_d = g\left(\sum_{k=1}^K x_{(d-1)+k} w_k\right)$$



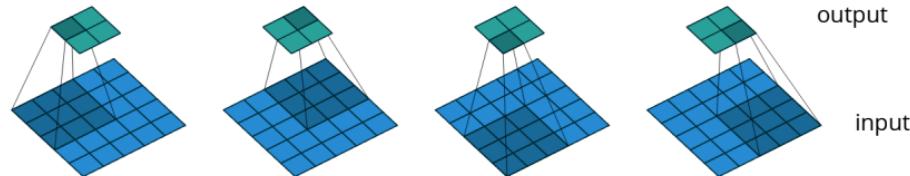
Strided convolution

the same idea extends to higher dimensions

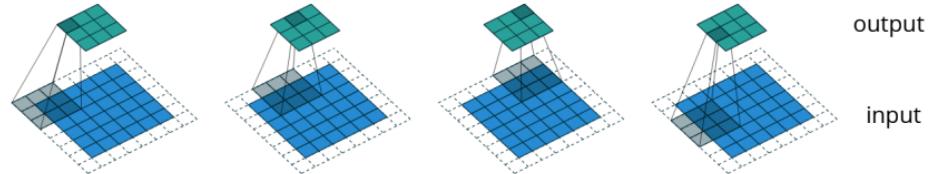
$$y_{d_1, d_2} = \sum_{k_1=1}^{K_1} \sum_{k_2=1}^{K_2} x_{p_1(d_1-1)+k_1, p_2(d_2-1)+k_2} w_{k_1, k_2}$$

different step-sizes for different dimensions

jump by 2



with padding



output length (for one dimension)

$$\left\lfloor \frac{D + \text{padding} - K}{\text{stride}} + 1 \right\rfloor$$

image: Dumoulin & Visin'16

Channels

so far we assumed a single input and output sequence or image

(Mx)
we have one $K_1 \times K_2$ filters per input-output channel combination $w \in \mathbb{R}^{M \times M' \times K_1 \times K_2}$

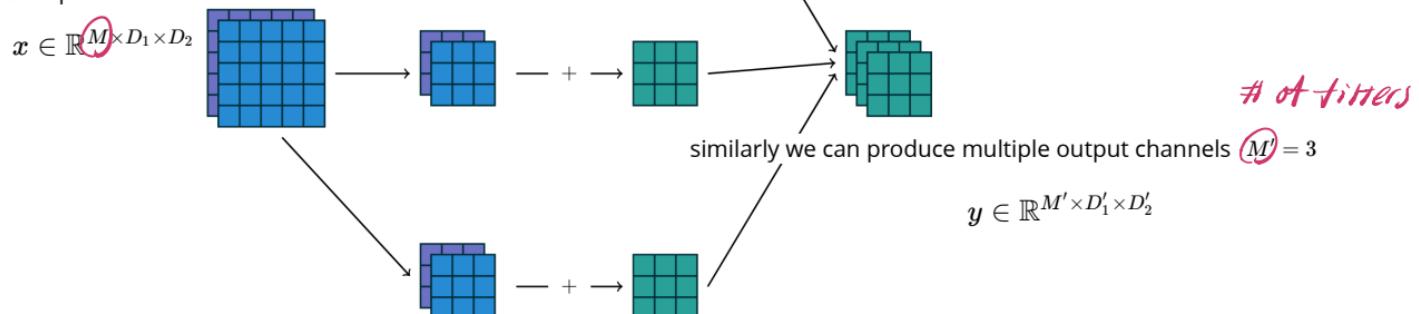
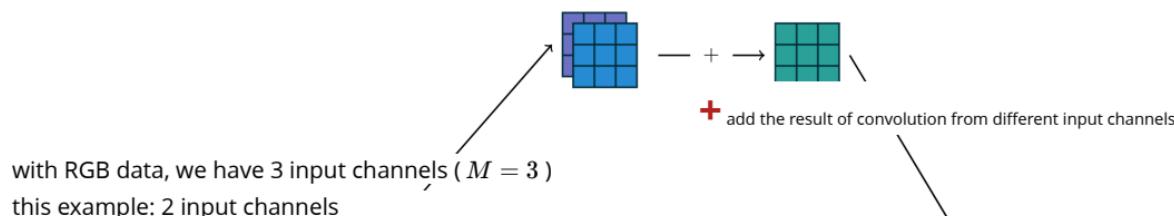


image: Dumoulin & Visin'16

Channels

so far we assumed a single input and output sequence or image

we can also add a *bias parameter* (b), one per each output channel $b \in \mathbb{R}^{M'}$

$$y_{m',d_1,d_2} = g\left(\sum_{m=1}^M \sum_{k_1} \sum_{k_2} w_{m,m',k_1,k_2} x_{m,d_1+k_1-1, d_2+k_2-1} + b_{m'}\right)$$

$$y \in \mathbb{R}^{M' \times D'_1 \times D'_2}$$

$$w \in \mathbb{R}^{M \times M' \times K_1 \times K_2}$$

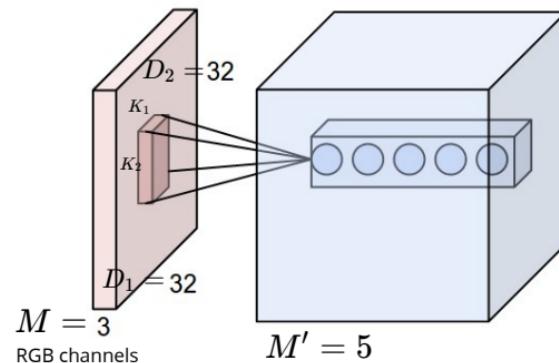


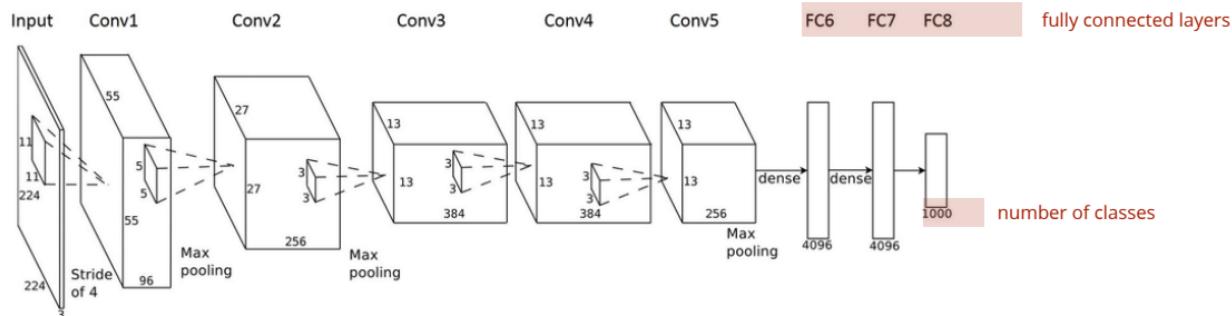
image: <https://cs231n.github.io/convolutional-networks/>

Convolutional Neural Network (CNN)

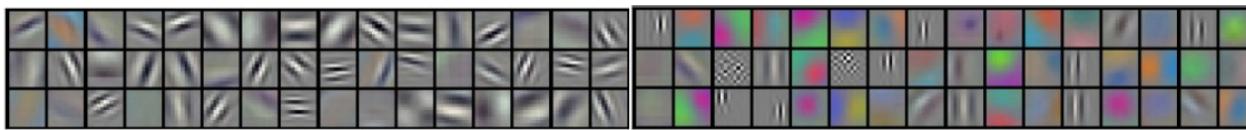
CNN or convnet is a neural network with convolutional layers (so it's a special type of MLP)

it could be applied to 1D sequence, 2D image or 3D volumetric data

example: conv-net architecture (derived from AlexNet) for image classification



↓ visualization of the convolution kernel at the first layer 11x11x3x96
96 filters, each one is 11x11x3. each of these is responsible for one of 96 feature maps in the second layer

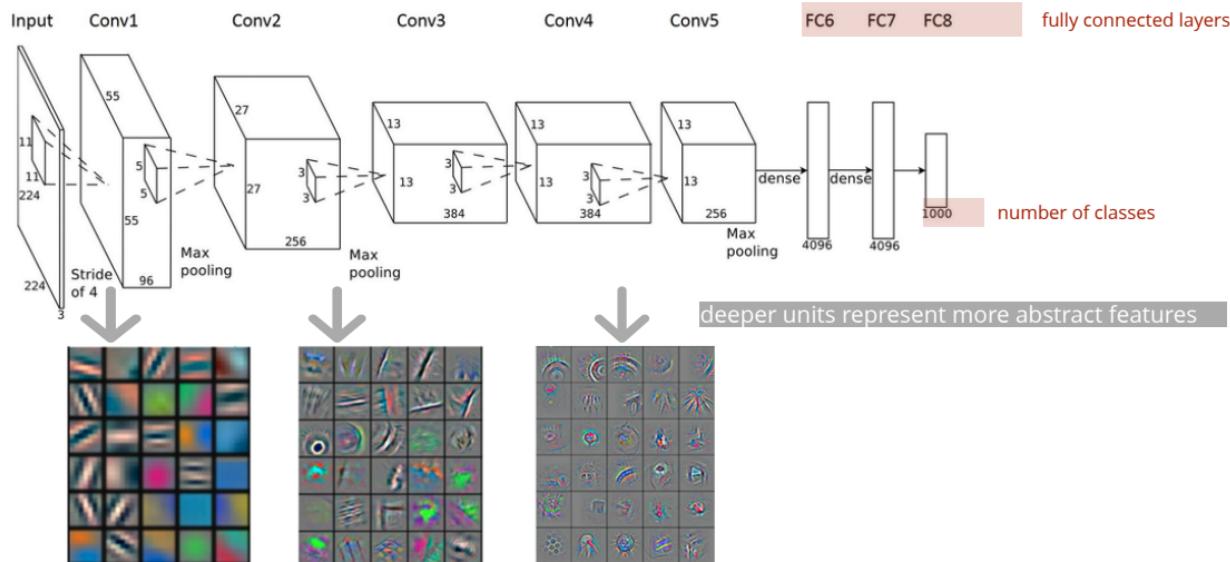


Convolutional Neural Network (CNN)

CNN or convnet is a neural network with convolutional layers (so it's a special type of MLP)

it could be applied to 1D sequence, 2D image or 3D volumetric data

example: conv-net architecture (derived from AlexNet) for image classification



Application: image classification

Convnets have achieved super-human performance in image classification

ImageNet challenge: > 1M images, 1000 classes

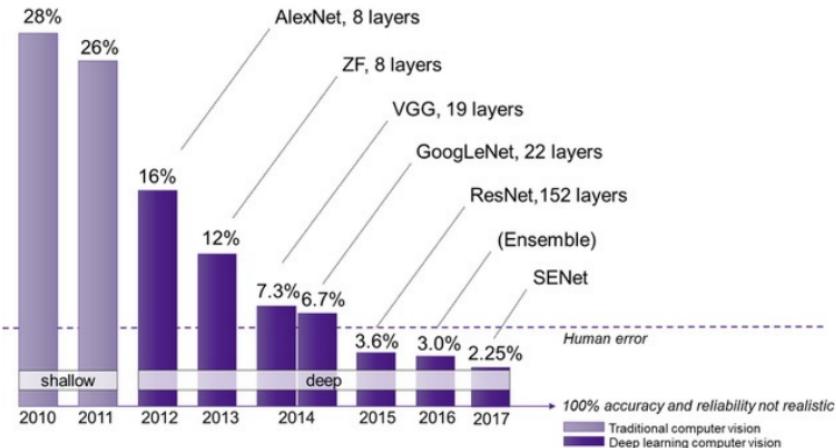
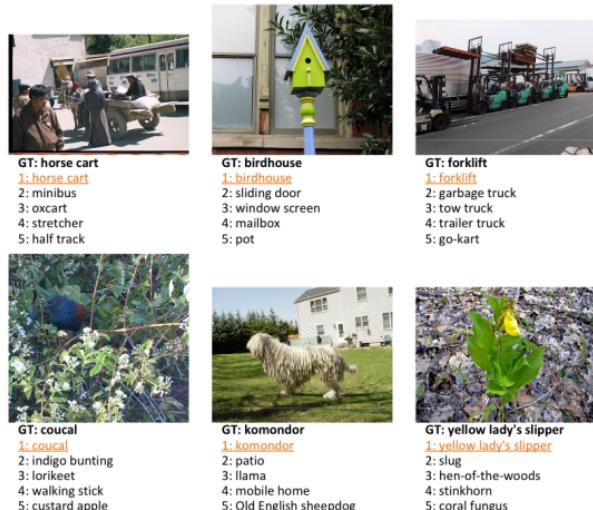


image credit: He et al'15, <https://semiengineering.com/new-vision-technologies-for-real-world-applications/>

Application: image classification

variety of increasingly deeper architectures have been proposed

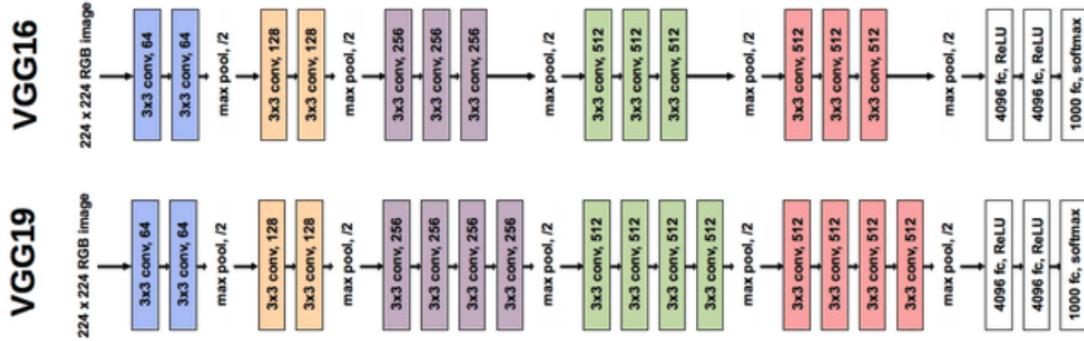


image credit: He et al'15, <https://semiengineering.com/new-vision-technologies-for-real-world-applications/>

Application: image classification

variety of increasingly deeper architectures have been proposed

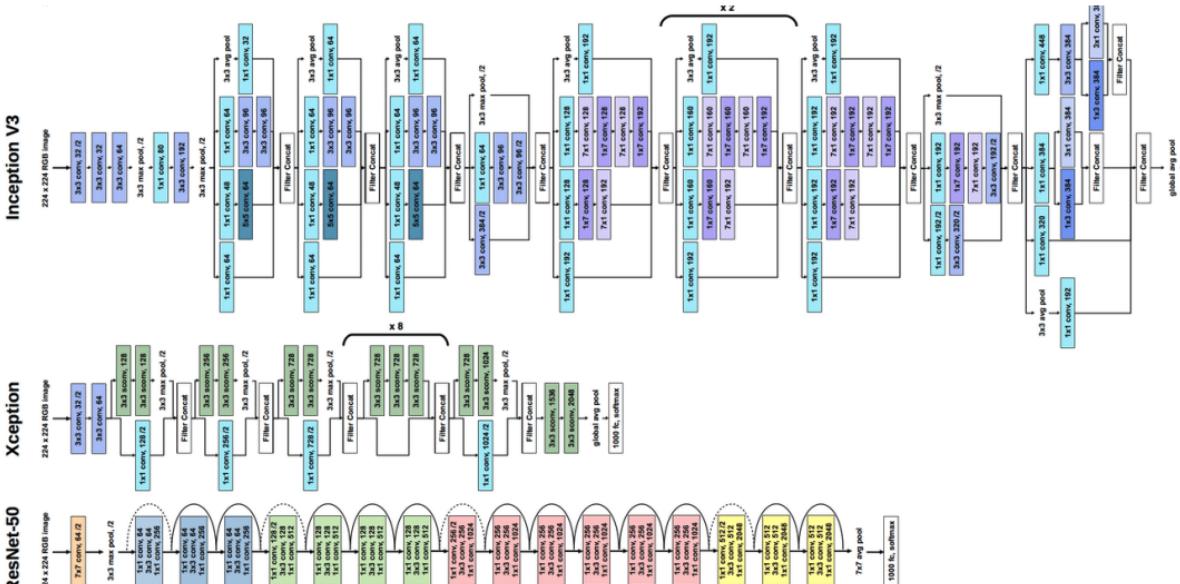


image credit: He et al'15, <https://semiengineering.com/new-vision-technologies-for-real-world-applications/>

Training: backpropagation through convolution

consider the strided 1D convolution op. $y_{m',d'} = \sum_m \sum_k w_{m,m',k} x_{m,p(d'-1)+k}$

output channel index input channel index filter index stride

using backprop. we have $\frac{\partial J}{\partial y_{m',d'}}$ so far and we need

1) $\frac{\partial y_{m',d'}}{\partial w_{m,m',k}}$ so as to get the gradients

$$\frac{\partial J}{\partial w_{m,m',k}} = \sum_{d'} \frac{\partial J}{\partial y_{m',d'}} \frac{\partial y_{m',d'}}{\partial w_{m,m',k}}$$



2) $\frac{\partial y_{m',d'}}{\partial x_{m,d}}$ to backpropagate to previous layer

$$\frac{\partial J}{\partial x_{d,m}} = \sum_{d',m'} \frac{\partial J}{\partial y_{m',d'}} \frac{\partial y_{m',d'}}{\partial x_{d,m}}$$



this operation is similar to multiplication by transpose of the parameter-sharing matrix (**transposed convolution**)

Naive implementation

consider the strided 1D convolution op. with stride 1. and single input-output channels

$$y_d = \sum_k w_k x_{d+k-1}$$

in practice most efficient implementation depends on the filter size (using FFT for large filters)



forward pass

```
1 def Conv1D(
2     x, # D (length)
3     w, # K (filter length)
4     ):
5
6     D, = x.shape
7     K, = w.shape
8     Dp = D - K + 1 #output length
9     y = np.zeros((Dp))
10    for dp in range(Dp):
11        y[dp] = np.sum(x[dp:dp+K] * w)
12    return y
```



backward pass

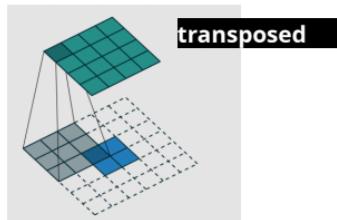
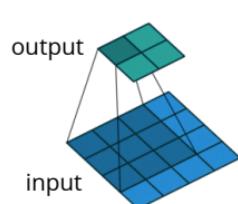
```
1 def Conv1DBackProp(
2     x, #D (length)
3     w, #K
4     dJdy,#Dp: error from layer above
5     ):
6
7     D, = x.shape
8     K, = w.shape
9     Dp, = dJdy.shape
10    dw = np.zeros_like(w)
11    dJdx = np.zeros_like(x)
12    for dp in range(Dp):
13        dw += np.sum(dJdy[dp] * x[dp:dp+K],
14                    dJdx[dp:dp+K] += dJdy[dp:dp+K] * w
15    return dJdx, dw #error to layer below and weight update
```

Transposed Convolution

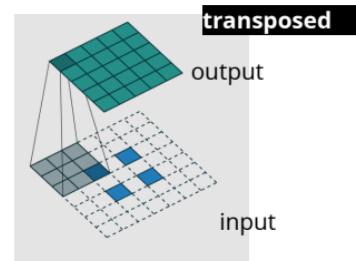
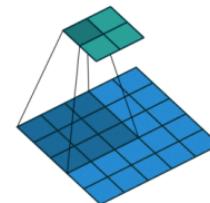
Transposed convolution (aka deconvolution) recovers the shape of the original input

Convolution with **no stride** and its transpose

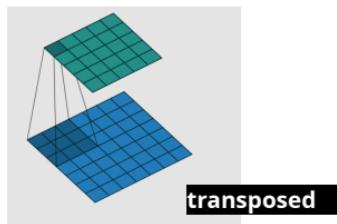
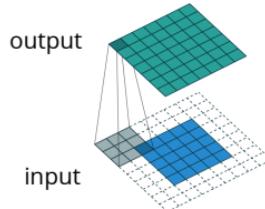
no padding of the original convolution corresponds to full padding of in transposed version



Convolution **with stride** and its transpose



full padding of the original convolution corresponds to no padding of in transposed version

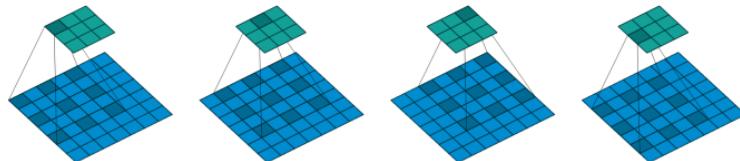


this can be used for up-sampling (opposite of stride/pooling)
as expected the transpose of a transposed convolution
is the original convolution

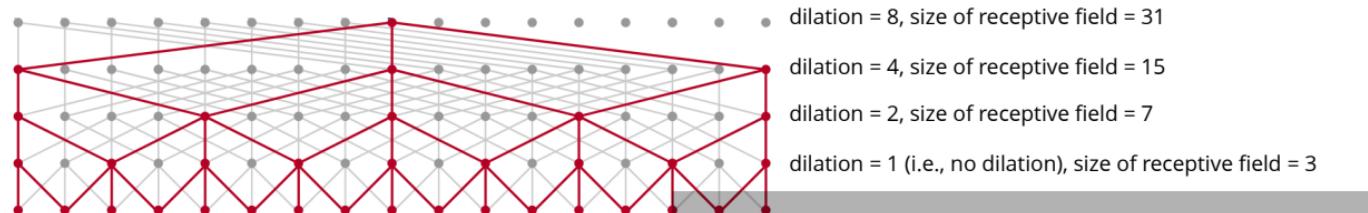
image: Dumoulin & Visin'16

Dilated Convolution

Dilated (aka atrous) convolution



this can be used to create exponentially large receptive field in few layers



in contrast to stride, dilation does not lose resolution

output length (for one dimension)

$$\left\lfloor \frac{D + \text{padding} - \text{dilation} \times (K-1) - 1}{\text{stride}} + 1 \right\rfloor$$



```
1 torch.nn.Conv2d(in_channels, out_channels, kernel_size,
                  stride=1, padding=0, dilation=1, groups=1, bias=True,
                  padding_mode='zeros')
```

image credits: Kalchbrenner et al'17, Dumoulin & Visin'16

Structured Prediction

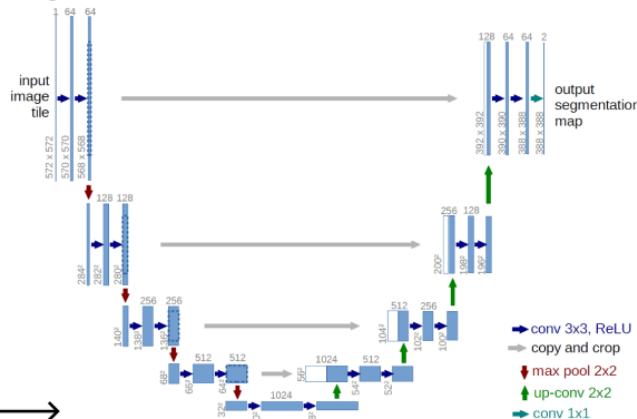
the output itself may have (image) structure (e.g., predicting text, audio, image)

example in (semantic) segmentation, we classify each pixel

loss is the sum of cross-entropy loss across the whole image



variety of architectures... one that performs well is **U-Net** →



transposed convolution (upconv), concatenation, and skip connection are common in architecture design
architecture search (i.e., combinatorial hyper-parameter search) is an expensive process and an active research area

image:https://sthalles.github.io/deep_segmentation_network/

Summary

convolution layer introduces an **inductive bias** to MLP

equivariance as an inductive bias:

- translation of the same model is applied to produce different outputs (pixels)
- the layer is equivariant to **translation**
- achieved through **parameter-sharing**

conv-nets use combinations of

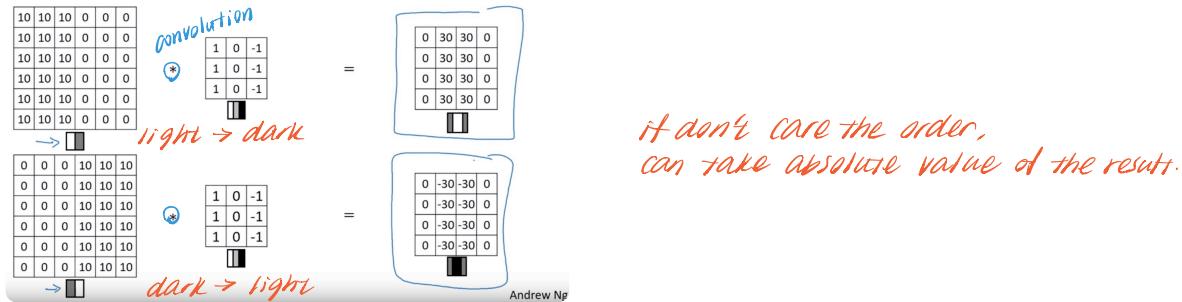
- convolution layers
- ReLU (or similar) activations
- pooling and/or stride for down-sampling
- skip-connection and/or batch-norm to help with optimization / regularization
- potentially fully connected layers in the end

training

- backpropagation (similar to MLP)
- SGD or its improved variations with adaptive learning rate
- monitor the validation error for early stopping

Edge detection

Vertical edge detection



Horizontal edge detection

$$\begin{array}{|c|c|c|c|c|c|} \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 10 & 10 & 10 \\ \hline 0 & 0 & 0 & 10 & 10 & 10 \\ \hline 0 & 0 & 0 & 10 & 10 & 10 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -1 & -1 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 0 & 0 & 0 & 0 \\ \hline 30 & 10 & -10 & -30 \\ \hline 30 & 10 & -10 & -30 \\ \hline 0 & 0 & 0 & 0 \\ \hline \end{array}$$

Notation

If layer l is a convolution layer:

$$\begin{aligned} f^{[l]} &= \text{filter size} & \text{Input: } & n_H^{[l-1]} \times n_W^{[l-1]} \times n_C^{[l-1]} \leftarrow \\ p^{[l]} &= \text{padding} & \text{Output: } & n_H^{[l]} \times n_W^{[l]} \times n_C^{[l]} \leftarrow \\ s^{[l]} &= \text{stride} & n_{HW}^{[l]} &= \left\lfloor \frac{n_H^{[l-1]} \times n_W^{[l-1]} - f^{[l]} + 2p^{[l]}}{s^{[l]}} + 1 \right\rfloor \\ n_c^{[l]} &= \text{number of filters} & \text{Activations: } & A^{[l]} \rightarrow n_H^{[l]} \times n_W^{[l]} \times n_C^{[l]} \\ \Rightarrow \text{Each filter is: } & f^{[l]} \times f^{[l]} \times n_C^{[l]} & \text{Weights: } & f^{[l]} \times f^{[l]} \times n_C^{[l-1]} \times n_C^{[l]} \\ & \text{bias: } n_C^{[l]} = (1, 1, 1, n_C^{[l]}) & \text{if bias in layer } l. & n_C^{[l]} \times n_H^{[l]} \times n_W^{[l]} \end{aligned}$$